

# Parallel Autoencoder

## High Performance Computing for Data Science

Silvanus Bordignon  
xxxxxx  
Univesity of Trento  
Trento, Italy  
maedje@typst.app

Ettore Sasgiorato  
247178  
University of Trento  
Trento, Italy  
ettore.saggiorato@studenti.unitn.it

**Abstract—TODO**

**Index Terms—HighPerformanceComputing, DeepLearning**

### I. INTRODUCTION

Aim of the project. Introduce that we use a small neural network, so the network won't be split between multiple nodes.

#### A. Instruction for Reproducibility and Building

cmake. runs on linux. Needs GCC and C++20.

### II. STATE OF THE ART / RELATED WORKS

### III. LIBRARIES AND DATASETS

#### A. Libraries

- stb
- Eigen
- Gtest
- GBench

#### B. Datasets

### IV. METHODOLOGY AND IMPLEMENTATION DETAILS

The structure of the operations of nn is fundamental, our objective is to maximize the TOPS.

- Critical thought about the cost of allocating memory and the usage of the heap/stack -> reference to the appendix
- Considering that Eigen can parallelize the workload thoughts about avoiding too much context switch on the CPU -> reference appendix
- Difference between when one has a GPU and doesn't in the dataloader
- Why we are doing distributed training instead of e.g. sharing the model between machines (Efficiency of course :). Reference a few papers plz

We planned to work sequentially:

- 1) Basic implementation single threaded
- 2) Activating Eigen's OpenMP parallelization
- 3) OpenMP the rest of hte code
- 4) MPI -> check how to use infiniband/omnipath
- 5) MPI + OpenMP

Say that methods are benchmarked individually to see where we gain the best performance and where it becomes worse.

#### A. Unit Testing

Performed to be sure that modules work correctly and that when parallelizing/using mpi nothing breaks.

#### B. Basic Implementation

- Why many parts of the NN are set as a template library.
- APIs inspired from PyTorch python's apis.

- 
- Dataloader
  - Linear layer
  - ReLU
  - Sigmoid
  - Encoder/Decoder
  - Loss
  - Backpropagation -> reference from the book (need to get it lol)
  - Gradient Descent

#### C. Eigen parallelization

#### D. OpenMP

#### E. MPI

> How we are parallelizing on multiple nodes > What we are sharing and Why > ABLATION on the distributed weights

#### F. Combo: MPI + OpenMP

### V. SYSTEM DESCRIPTION

### VI. EXPERIMENTS

#### A. Evaluation

What we evaluate:

- Speedup
- Efficiency
- Scalability
  - Strong
  - Weak

a) *What we expect:*

#### B. TODO: experiment combos and their results compared.

What we can learn from this, how it goes.

## VII. TABLES AND DATA

## VIII. ADDITIONAL LISTINGS

## REFERENCES

## APPENDIX A.

“asd”, adwaoids

## APPENDIX B. IF NEEDED WE CAN ADD A TITLE TO THE APPENDIX :)

- Ablations about where performances explode + some thoughts about the design process
- heap vs stack with Eigen (`Matrix<float, ...>` vs `MatrixXf`)
- using vs not using eigen parallelization: it's parallelization (single loop) on a set of data vs a possible openmp done by us where data is worked on in parallel